

# Docker Introduction



docker®

By Eng: Abdelrahman Ayman Lotfy

DevOps track

## Table of contents

Why docker .....	3
History of docker .....	3
Docker build process .....	5
Internal process .....	5
Docker Run process .....	6
Internal process .....	7
Docker Instructions .....	8
Docker Commands .....	11
Conclusion .....	12

A large, light blue watermark of the Docker logo is centered on the page. It features a stylized ship with a stack of five containers on its deck, sailing on a wave.

docker®

## why Docker?

- **Solve the "it works on my machine" problem:** Docker addressed inconsistencies in application behavior across different environments due to varying dependencies, libraries, and configurations.
  - **Simplify containerization:** Built on Linux container technologies like LXC, Docker provided a user-friendly API and ecosystem for building, shipping, and managing containers.
  - **Enable portability:** Containers package applications with all required components, ensuring consistent execution across any Docker-supported system.
  - **Improve efficiency:** Offers lightweight, isolated environments using OS-level virtualization, reducing the overhead of full virtual machines.
  - **Automate deployment:** Streamlines the process of deploying applications with reproducible environments.
- 

## History of Docker

- **2008/2010:** DotCloud, a PaaS company, founded by Solomon Hykes, Kamel Founadi, and Sebastien Pahl in Paris, part of Y Combinator's Summer 2010 incubator.
- **2011:** DotCloud launched its PaaS platform; Docker began as an internal project to leverage container technology for application isolation.
- **March 2013:** Docker debuted at PyCon in Santa Clara, CA, and was released as open-source software, initially using LXC as its execution environment.
- **2013:** DotCloud rebranded to Docker Inc., focusing on containerization.

- **2014:**
  - ❖ Docker version 0.9 replaced **LXC** with libcontainer (written in Go) for better control and portability.
  - ❖ Partnerships formed with Red Hat (Fedora, OpenShift integration), Microsoft (Windows Server support), Amazon EC2, and IBM.
  - ❖ Docker Compose beta released in December 2013, with version 1.0 in October 2014.
- **2015:** New open standard for containers announced with multiple companies.
- **2016:** Swarm mode integrated into Docker Engine (version 1.12).
- **2017:** Moby project launched for open R&D on container systems; Docker usage grew 160% on LinkedIn by January.
- **2019:** Enhanced Windows support via WSL 2 announced in May.
- **2021:** Docker introduced a Personal Plan and ended free Docker Desktop for large businesses in August.
- **2023:** Docker acquired AtomicJar in December to enhance testing features.
- **2024:** Docker celebrated its 11th anniversary, solidifying its role in the container revolution.
- **Present:** Docker Engine remains open-source, widely used in cloud-native ecosystems, with ongoing development by **Docker Inc.**

## Docker Build Process

- **Parse Dockerfile:** Docker reads the Dockerfile in the current directory (or specified path) to understand the instructions for building the image.
- **Create a build context:** Collects all files in the current directory (or specified context) to be used during the build, excluding those listed in .dockerignore.
- **Set up temporary container:** Docker creates an intermediate container for each instruction in the Dockerfile to execute commands in isolation.
- **Execute Dockerfile instructions:**
  - ❖ Processes commands like FROM (sets base image), COPY/ADD (adds files to image), RUN (executes commands), ENV (sets environment variables), and others.
  - ❖ Each instruction creates a new layer in the image, cached for reuse in future builds unless changes occur.
- **Generate image layers:** Each instruction's result is stored as a read-only layer, forming the final image's filesystem.
- **Tag the image:** Assigns a tag (e.g., myapp:latest) to the built image, as specified by the -t flag or default naming.
- **Save the image:** Stores the final image in the local Docker image registry, ready for use or pushing to a remote registry (e.g., Docker Hub).
- **Clean up:** Removes intermediate containers used during the build to save space, unless --no-cache is used.

### Internal process:

- ❖ **Client-Server communication:** The Docker CLI sends the build request to the Docker daemon (via REST API over a Unix socket or TCP).

- ❖ **BuildKit integration:** If enabled, BuildKit (Docker's build backend) optimizes the build with parallel processing and efficient caching.
- ❖ **Storage driver:** The daemon uses a storage driver (e.g., overlay2, aufs) to manage image layers and copy-on-write filesystem operations.
- ❖ **Namespace isolation:** Leverages Linux namespaces (e.g., PID, mount, network) for container isolation during intermediate container execution.
- ❖ **Cgroups:** Applies control groups to limit resources (CPU, memory) for build processes.
- ❖ **Image manifest:** Creates a JSON manifest detailing the image's layers, configuration, and metadata, stored in the local registry.

## Docker Run Process

- **Locate the image:** Docker checks the local registry for the specified image; if not found, it pulls it from a remote registry (e.g., Docker Hub) unless `--pull=never` is set.
- **Create a container:** Allocates a new container from the specified image, setting up an isolated environment with its own filesystem, network, and process space.
- **Configure container settings:**
  - ❖ Applies options from the docker run command, such as port mappings (-p), environment variables (-e), volumes (-v), or detach mode (-d).
  - ❖ Sets the container's entrypoint and command as defined in the Dockerfile (ENTRYPOINT and CMD) or overridden via command-line arguments.
- **Start the container:** Initializes the container's runtime, executing the specified command or entrypoint in the isolated environment.

- **Allocate resources:** Assigns CPU, memory, and other resources based on defaults or user-specified limits (e.g., `--memory`, `--cpu-shares`).
- **Set up networking:** Configures the container's network, connecting it to the specified network (e.g., bridge, host) and mapping ports if requested.
- **Run the application:** Executes the container's main process, outputting logs to the terminal (unless detached with `-d`) and keeping the container running until the process exits.
- **Handle container lifecycle:** If the process exits, the container stops unless configured to restart (e.g., `--restart=always`); stopped containers remain for inspection unless removed with `--rm`.

### Internal process:

- ❖ **Client-Server interaction:** The Docker CLI sends the run request to the Docker daemon via the API.
- ❖ **Container runtime:** Uses a runtime (e.g., runc, containerd) to create and manage the container, interfacing with the Linux kernel.
- ❖ **Namespaces:** Sets up Linux namespaces (PID, mount, network, user, UTS, IPC) to isolate the container's processes, filesystem, hostname, and network stack.
- ❖ **Cgroups:** Configures control groups to enforce resource limits and track resource usage.
- ❖ **Storage driver:** Mounts the image's read-only layers with a writable layer (using storage drivers like overlay2) for the container's filesystem.

- ❖ **Networking setup:** Configures network namespaces, virtual ethernet pairs (veth), and bridges (e.g., docker0) for connectivity, with iptables rules for port mapping.
- ❖ **Seccomp and AppArmor:** Applies security profiles (seccomp filters, AppArmor policies) to restrict system calls and enhance container security.
- ❖ **Log management:** Streams container logs to the CLI or a logging driver (e.g., json-file, syslog) based on configuration.

## Docker Instructions

All the following inside a dockerfile:

### 1. FROM

Defines the **base image**.

FROM ubuntu:20.04

---

### 2. RUN

Runs a command during the image build (creates a new layer).

RUN apt-get update && apt-get install -y python3

---

### 3. CMD

Defines the **default command** to run when a container starts.  
(Only one CMD, last one overrides).

CMD ["python3", "app.py"]



#### 4. ENTRYPOINT

Similar to CMD, but **cannot be overridden** easily. Used for fixed executables.

```
ENTRYPOINT ["python3", "app.py"]
```

---

#### 5. COPY

Copies files from the **host machine** → **image**.

```
COPY ./app
```

---

#### 6. ADD

Like COPY, but supports remote URLs and tar extraction.

```
ADD https://example.com/file.tar.gz /app/
```

---

#### 7. WORKDIR

Sets the **working directory** inside the container.

```
WORKDIR /app
```

---

#### 8. ENV

Set environment variables.

```
ENV PORT=5000
```

---

#### 9. EXPOSE

Documents the port your app will use (doesn't publish it automatically).

```
EXPOSE 5000
```

## 10. VOLUME

Creates a mount point for external volumes.

**VOLUME ["/data"]**

---

## 11. ARG

Defines build-time variables (used only during docker build).

**ARG APP\_VERSION=1.0**

---

## 12. LABEL

Add metadata to the image.

**LABEL maintainer="your email"**

**LABEL version="1.0"**

---

## 13. USER

Set the user that the container will run as.

**USER appuser**

---

## 14. SHELL

Change the default shell used by RUN commands.

**SHELL ["/bin/bash", "-c"]**

# Docker Commands

## Image Management

<code>docker build -t name:tag .</code>	# Build an image from Dockerfile
<code>docker images</code>	# List all images
<code>docker rmi image_id</code>	# Remove an image

## Container Lifecycle

<code>docker run image</code>	# Create & start a container
<code>docker run -it ubuntu bash</code>	# Run container with interactive terminal
<code>docker ps</code>	# List running containers
<code>docker ps -a</code>	# List all containers
<code>docker stop container_id</code>	# Stop a container
<code>docker start container_id</code>	# Start a container
<code>docker rm container_id</code>	# Remove a container

## Debugging & Logs

<code>docker logs container_id</code>	# Show container logs
<code>docker exec -it container_id bash</code>	# Access container shell
<code>docker inspect container_id</code>	# Detailed info about container

## Networking

<code>docker port container_id</code>	# Show port mappings
<code>docker network ls</code>	# List networks
<code>docker network create mynet</code>	# Create a custom network

## Volumes (Persistent Data)

<code>docker volume ls</code>	# List volumes
<code>docker volume create myvol</code>	# Create a volume
<code>docker run -v myvol:/data image</code>	# Mount volume inside container

## Docker Hub (Registry)

<code>docker login</code>	# Login to Docker Hub
<code>docker pull image</code>	# Download image
<code>docker push image</code>	# Upload image

---

## Conclusion

- Docker solves real-world problems like environment inconsistencies, dependency conflicts, and heavy virtual machines.
- It became popular because it is lightweight, portable, and fast.
- Images are the blueprints, and containers are the running instances of those images.
- `docker build` creates an image (read-only layers), while `docker run` creates a container (with a writable layer).
- Dockerfile instructions (`FROM`, `RUN`, `COPY`, `CMD`, `ENTRYPOINT`, etc.) define how to build images.
- Docker has become a standard in DevOps and cloud computing, and is often combined with Kubernetes for orchestration.

Docker makes building, shipping, and running applications easier, faster, and more reliable everywhere.